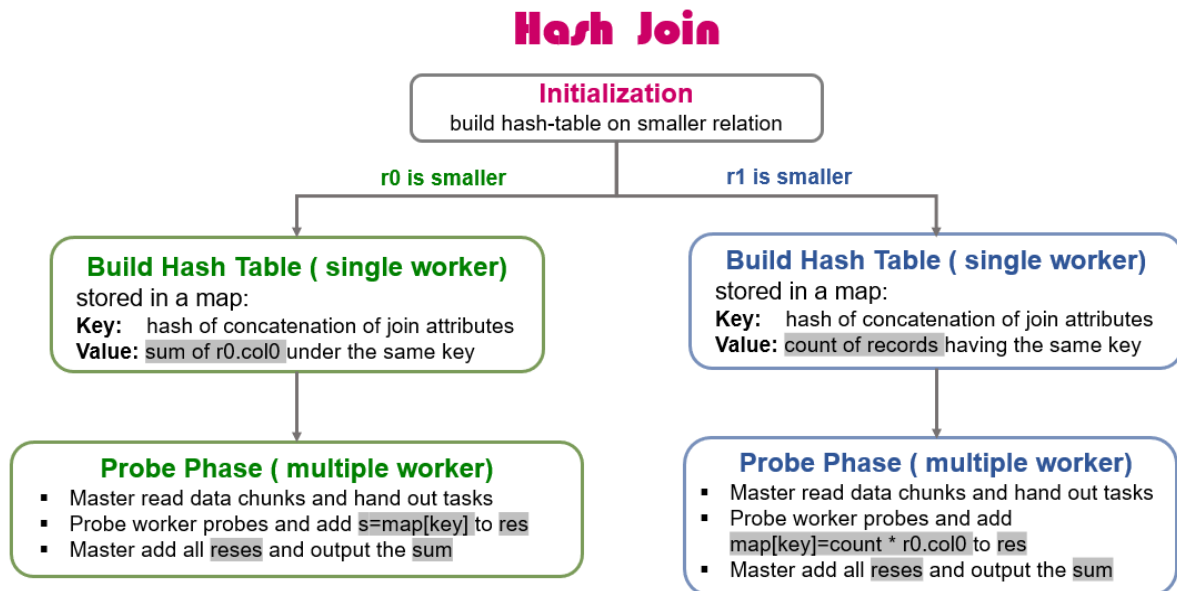


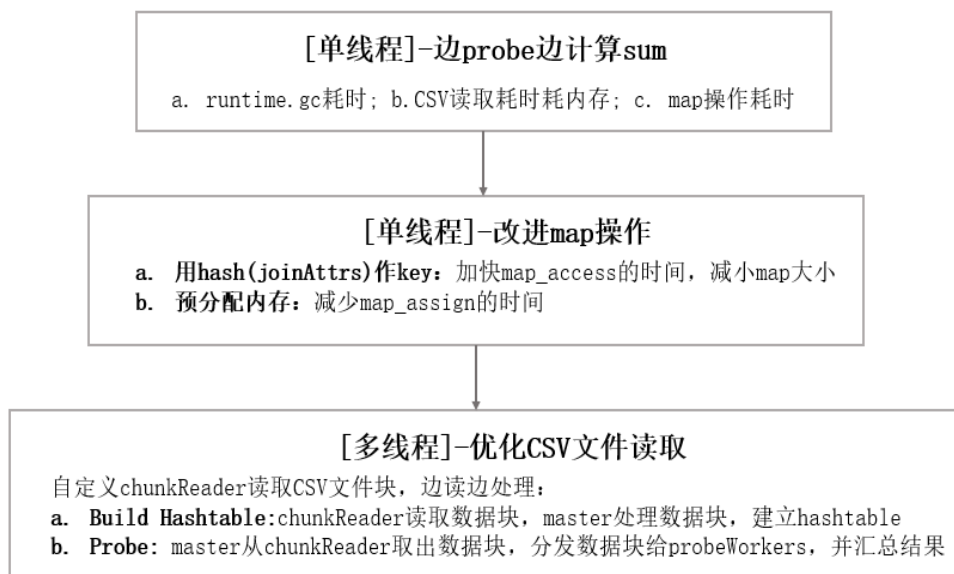
1. Ideas

考虑到MergeJoin要求排序，是全有序的，而HashJoin是先分类，缩小范围后再join。所以在Join Attribute无序的情况下，MergeJoin是要比HashJoin慢的。看了一下t文件夹下要排序的文件，第一列是依次递增的，之后的列是无序的。不知道之后正式测试的文件会是什么样，这里还是只考虑了HashJoin。基本思路如下图所示。



2. PProf

基本思路如下图所示。



Phase 1: [单线程]--边probe边计算sum

这一步相对于Example的优化主要体现在：

1. 减少了map的大小。原来在相同key下的map需要存相应row number的slice，现在只需要存这些row的col0的和。

2. 减少了probe后再用row number去提取相应col0求和的开销。

但是从pprof的结果来看，提升的空间还有很大：a. GC耗时太长；b. CSV读取操作耗时耗内存严重；c. map_assign和map_access耗时太长，map扩容重映射的情况也很严重。pprof CPU和pprof mem的结果分别如下图所示。

```
(pprof) top20 -cum
Showing nodes accounting for 8.08s, 47.42% of 17.04s total
Dropped 126 nodes (cum <= 0.09s)
Showing top 20 nodes out of 72
flat flat% sum% cum cum% runtime.systemstack
0.01s 0.059% 0.059% 9.17s 53.81% main.main
0 0% 0.059% 8.40s 49.30% runtime.main
0 0% 0.059% 8.40s 49.30% runtime.main
0 0% 0.059% 8.36s 49.06% main.Join
0 0% 0.059% 8.33s 48.88% runtime.gcBgMarkWorker.func2
0.24s 1.41% 1.47% 8.33s 48.88% runtime.gcDrain
0 0% 1.47% 8.30s 48.71% runtime.gcBgMarkWorker
2.10s 12.32% 13.79% 7.21s 42.31% runtime.scanobject
0.11s 0.65% 14.44% 4.45s 26.12% main._readCSVFileIntoTbl
0.06s 0.35% 14.79% 3.08s 18.08% encoding/csv.(*Reader).Read
0.65s 3.81% 18.60% 3.02s 17.72% encoding/csv.(*Reader).readRecord
0.17s 1% 19.60% 2.74s 16.08% main._buildHashTable
1.91s 11.21% 30.81% 2.67s 15.67% runtime.findObject
1.05s 6.16% 36.97% 2.30s 13.50% runtime.mapassign_faststr
1.12s 6.57% 43.54% 2.14s 12.56% runtime.greyobject
0.26s 1.53% 45.07% 1.33s 7.81% runtime.mallocgc
0 0% 45.07% 1.24s 7.28% runtime.growslice
0.05s 0.29% 45.36% 1.17s 6.87% main._probe
0.35s 2.05% 47.42% 1.10s 6.46% runtime.evacuate_faststr
0 0% 47.42% 1.10s 6.46% runtime.growWork_faststr

(pprof) list main._buildHashTable
Total: 17.04s
ROUTINE ===== main._buildHashTable in D:\Coding\Golang\src\fftest\Full-PingCAP\tidb\join\join.go
170ms 2.74s (flat, cum) 16.08% of Total
47: }
48:
49: func _buildHashTable(data [][]string, offset []int) map[string]uint64 {
50:     var keyBuffer []byte
51:     hashtable := make(map[string]uint64)
52:     for _, row := range data {
53:         for _, off := range offset {
54:             keyBuffer = append(keyBuffer, []byte(row[off])...)
55:         }
56:         v, err := strconv.ParseUint(row[0], 10, 64)
57:         if err != nil {
58:             panic("JoinExample panic\n" + err.Error())
59:         }
60:         hashtable[string(keyBuffer)] += v
61:         keyBuffer = keyBuffer[:0]
62:     }
63:     return hashtable
64: }
65: }

(pprof) top20 -cum
Showing nodes accounting for 2993.49MB, 100% of 2993.99MB total
Dropped 3 nodes (cum <= 14.97MB)
flat flat% sum% cum cum% main.Join
0 0% 0% 2993.49MB 100% main.main
0 0% 0% 2993.49MB 100% runtime.main
1415.04MB 47.26% 47.26% 2445.59MB 81.68% main._readCSVFileIntoTbl
1030.55MB 34.42% 81.68% 1030.55MB 34.42% encoding/csv.(*Reader).Read
547.90MB 18.30% 100% 547.90MB 18.30% main._buildHashTable

(pprof) list main._buildHashTable
Total: 2.92GB
ROUTINE ===== main._buildHashTable in D:\Coding\Golang\src\fftest\Full-PingCAP\tidb\join\join.go
547.90MB 547.90MB (flat, cum) 18.30% of Total
55: }
56: v, err := strconv.ParseUint(row[0], 10, 64)
57: if err != nil {
58:     panic("JoinExample panic\n" + err.Error())
59: }
547.90MB 547.90MB 60: hashtable[string(keyBuffer)] += v
61:
62: keyBuffer = keyBuffer[:0]
63: }
64: return hashtable
65: }

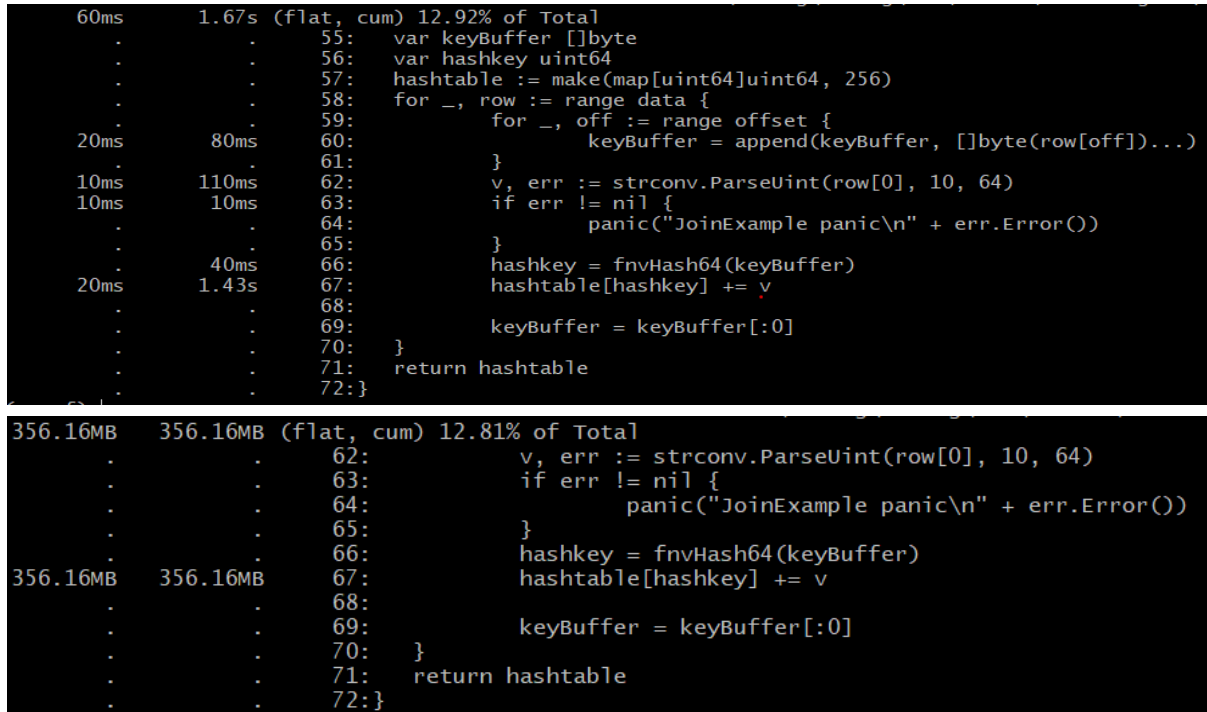
(pprof) list main._readCSVFileIntoTbl
Total: 2.92GB
ROUTINE ===== main._readCSVFileIntoTbl in D:\Coding\Golang\src\fftest\Full-PingCAP\tidb\join\join.go
1.38GB 2.39GB (flat, cum) 81.68% of Total
32: }
33: defer csvFile.Close()
34:
35: csvReader := csv.NewReader(csvFile)
36: for {
37:     row, err := csvReader.Read()
38:     if err == io.EOF {
39:         break
40:     } else if err != nil {
41:         panic("ReadFileIntoTbl " + f + " fail\n" + err.Error())
42:     }
1.38GB 1.38GB 43: tbl = append(tbl, row)
44: }
45:
46: return tbl
47: }
48: }
```

Phase 2: [单线程]--优化map操作

这一步相对于Example的优化主要体现在：

1. 用hash(joinAttrs) uint64类型, 替代joinAttrs string类型, 作为map key以加速map_access, 同时减小map的大小
2. 预分配map空间, 以减少map扩容重映射的时间, 减少map_assign的时间

pprof CPU和pprof mem的结果分别如下图所示, 可以看到优化后的map耗时减少了1秒左右, 内存占用也减少了200M左右。



Phase 3: [多线程]--优化CSV文件读取

从之前的pprof结果可以看出: 使用csv包一行一行地把数据读到一个slice中的方式, 在时间上csv/decoding十分耗时, 在空间上存放所有数据的slice开销巨大。于是决定放弃使用csv包, 同时为了节省时间, 采用读数据和处理数据同步进行的方式。程序中一共有三个角色: chunkReader、master、worker。其中chunkReader负责读取数据块, 当chunkReader读完一个数据块后会通知master有新的数据块, master收到通知后, 再去把新的数据块取过来, 然后告诉chunkReader可以读下一个数据块了, 并同时新的数据块封装成task分发给worker处理。worker处理完自己的数据块后, 将结果返还给master, 由master汇总输出。处理思路大致如下:

1. **Build Hashtable.** chunkReader负责读数据块, master接收数据块并且处理数据块, 建立hashtable。为了内存复用, master有一个dataChunk缓冲区, 每次都将从chunkReader取下来的chunk放到这里
2. **Probe.** master负责从chunkReader那里取chunk, 并将chunk分发给probeWorker。为了内存复用, master有一个[]dataChunk的缓冲区, 在分发task给probeWorker时, 会在task中指明该task对应的chunk在[]dataChunk中的index, 当probeWorker处理完task后, 会将这个index和处理结果一起返还给master, 告诉master dataChunk[index]这个chunk可用了。如果chunkReader读数据比较快, 而[]dataChunk中没有可用的chunk了, master才会增大[]dataChunk的空间。

pprof CPU和pprof mem的结果分别如下图所示。可以看出, 时间上相比于example优化了4倍左右, 空间上优化了10倍左右。但最后map操作还是成为了系统的瓶颈。

```

Duration: 2.45s, Total samples = 5.61s (228.91%)
Entering interactive mode (type "help" for commands, "o" for options)
(pprof) top20 -cum
Showing nodes accounting for 5150ms, 91.80% of 5610ms total
Dropped 34 nodes (cum <= 28.05ms)
Showing top 20 nodes out of 30

```

flat	flat%	sum%	cum	cum%	
2110ms	37.61%	37.61%	3320ms	59.18%	main.(*hashJoin).probeworker1
0	0%	37.61%	2060ms	36.72%	main.main
0	0%	37.61%	2060ms	36.72%	runtime.main
0	0%	37.61%	2050ms	36.54%	main.(*hashJoin).Join
0	0%	37.61%	2050ms	36.54%	main.Join
0	0%	37.61%	1650ms	29.41%	main.(*hashJoin).BuildHashtable
630ms	11.23%	48.84%	1650ms	29.41%	main.(*hashJoin).buildHashtable1
1070ms	19.07%	67.91%	1140ms	20.32%	runtime.mapaccess2_fast64
580ms	10.34%	78.25%	810ms	14.44%	runtime.mapassign_fast64
550ms	9.80%	88.06%	550ms	9.80%	main.(*chunkReader).WriteChunk
0	0%	88.06%	400ms	7.13%	main.(*hashJoin).Probe
0	0%	88.06%	210ms	3.74%	main.(*chunkReader).RunReading
60ms	1.07%	89.13%	170ms	3.03%	runtime.evacuate_fast64
0	0%	89.13%	170ms	3.03%	runtime.growWork_fast64
0	0%	89.13%	160ms	2.85%	bufio.(*Reader).Read
0	0%	89.13%	160ms	2.85%	io.ReadAtLeast
0	0%	89.13%	160ms	2.85%	io.ReadFull
150ms	2.67%	91.80%	150ms	2.67%	runtime.memmove
0	0%	91.80%	110ms	1.96%	internal/poll.(*FD).Read
0	0%	91.80%	110ms	1.96%	os.(*File).Read

```

(pprof) top20 -cum
Showing nodes accounting for 216.39MB, 100% of 216.39MB total

```

flat	flat%	sum%	cum	cum%	
0	0%	0%	208.37MB	96.29%	main.main
0	0%	0%	208.37MB	96.29%	runtime.main
0	0%	0%	206.65MB	95.50%	main.Join
0	0%	0%	198.65MB	91.80%	main.(*hashJoin).Join
0	0%	0%	166.58MB	76.98%	main.(*hashJoin).BuildHashtable
166.58MB	76.98%	76.98%	166.58MB	76.98%	main.(*hashJoin).buildHashtable1
32.07MB	14.82%	91.80%	32.07MB	14.82%	main.(*hashJoin).Probe
8.02MB	3.71%	95.51%	8.02MB	3.71%	bufio.NewReaderSize
0	0%	95.51%	8.02MB	3.71%	main.(*chunkReader).RunReading
0	0%	95.51%	8MB	3.70%	main.(*hashJoin).init
8MB	3.70%	99.20%	8MB	3.70%	main.NewChunkReader
0	0%	99.20%	8MB	3.70%	main.NewHashJoin
1.72MB	0.8%	100%	1.72MB	0.8%	runtime/pprof.StartCPUProfile

```

240ms      280ms      384:      }
.          1.14s      385:      if count, ok := hj.hashtable[hashkey]; ok {
130ms      130ms      386:          res += val * count
.          .          387:      }

.          .          331:      }
.          .          332:      // store the count
162.57MB   162.57MB  333:      hj.hashtable[hashkey]++
.          .          334:
.          .          335:      attrShift = attrShift[:0]
.          .          336:      keyBuffer = keyBuffer[:0]
.          .          337:      attrShift = append(attrShift, i)
.          .          338:      }

```